

UNIVERSITY OF PENNSYLVANIA
ESE 650: LEARNING IN ROBOTICS
SPRING 2021
[03/18] HOMEWORK 3
DUE: 04/08 THU 11.59 PM ET

Changelog: This space will be used to note down updates/errata to the homework problems.

Instructions

Read the following instructions carefully before beginning to work on the homework.

- You will submit solutions typeset in \LaTeX on Gradescope (strongly encouraged). You can use `hw_template.tex` on Canvas in the “Homeworks” folder to do so. If your handwriting is *unambiguously legible*, you can submit PDF scans/tablet-created PDFs.
- Please start a new problem on a fresh page and mark all the pages corresponding to each problem. Failure to do so may result in your work not graded completely.
- Clearly indicate the name and Penn email ID of all your collaborators on your submitted solutions.
- **For each problem in the homework, you should mention the total amount of time you spent on it. This helps us gauge the perceived difficulty of the problems.**
- You can be informal while typesetting the solutions, e.g., if you want to draw a picture feel free to draw it on paper clearly, click a picture and include it in your solution. Do not spend undue time on typesetting solutions.
- You will see an entry of the form “HW 3 PDF” where you will upload the PDF of your solutions. You will also see entries like “HW 3 Problem 1 Code” where you will upload your solution for the respective problems. **For each programming problem, you should create a fresh Python file.** This file should contain all the code to reproduce the results of the problem and you will upload the .py file to Gradescope. If we have installed Autograder for a particular problem, you will use the Autograder.
- **You should include all the relevant plots in the PDF, without doing so you will not get full credit.** You can, for instance, export your Jupyter


notebook as a PDF (you can also use text cells to write your solutions) and export the same notebook as a Python file to upload your code.

- **Your PDF solutions should be completely self-contained. We will run the Python file to check if your solution reproduces the results in the PDF.**

Credit The points for the problems add up to 120. You only need to solve for 100 points to get full credit, i.e., your final score will be $\min(\text{your total points}, 100)$.

1 **Problem 1 (Policy Iteration, 40 points).** Consider the following Markov Decision
2 Process. The state-space is a 10×10 grid, cells that are obstacles are marked in gray.
3 The initial state of the robot is in blue and our desired terminal state is in green.
4 The robot gets a *reward* of 10 if it reaches the desired terminal state with a discount
5 factor of 0.9. At each non-obstacle cell, the robot can attempt to move to any of the
6 immediate neighboring cells using one of the four controls (North, East, West and
7 South). The robot cannot diagonally. The move succeeds with probability 0.7 and
8 with remainder probability 0.3 the robot can end up at some other cell as follows:

$$\begin{aligned} P(\text{moves north} \mid \text{control is north}) &= 0.7, \\ P(\text{moves west} \mid \text{control is north}) &= 0.1, \\ P(\text{moves east} \mid \text{control is north}) &= 0.1, \\ P(\text{does not move} \mid \text{control is north}) &= 0.1. \end{aligned}$$



9


10 Similarly, if the robot desired to go east, it may end up in the cells to its north,
11 south, or stay put at the original cell with total probability 0.3 and actually move
12 to the cell east with probability 0.7. The cost pays a cost of 1 (i.e., reward is -1)
13 for each control input it takes, regardless of the outcome. If the robot ends up at a
14 state marked as an obstacle, it gets a reward of -10 for each time-step that it remains
15 inside the obstacle cell.

16 We would like to implement policy iteration to find the best trajectory for the
17 robot to go from the blue cell to the green cell.

- 18 (a) **(10 points)** Carefully code up the above environment to run policy iteration.
19 You will need to think about how to code up the probability transition matrix
20 $\mathbb{R}^{100 \times 100} \ni T_{x,x'}(u) = P(x' \mid x, u)$, the run-time cost $q(x, u)$, and the
21 terminal cost $q_f(x)$. Policy iteration is easy to implement if you represent
22 all the above quantities as matrices and vectors. Plot the environment to
23 check if it confirms to the above picture.
24 (b) **(15 points)** Initialize policy iteration with a feedback control $u^{(0)}(x)$ where
25 the robot always goes east, this results in a policy $\pi^{(0)} = (u^{(0)}(\cdot), u^{(0)}(\cdot), \dots)$.
26 Write the code for policy evaluation to obtain the cost-to-go from every cell

- 1 in the above picture for this initial policy. Plot the value function $J^{\pi^{(0)}}(x)$
2 as a heatmap in the above picture.
- 3 (c) **(15 points)** Execute the policy iteration algorithm, you will iteratively per-
4 form policy evaluation and policy improvement steps. For the first 4 itera-
5 tions, plot the feedback control $u^{(k)}(x)$ (using arrows as shown in the lecture
6 notes (https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.arrow.html,
7 you can also write the control input in the cell). You should color the cell
8 using the value function $J^{\pi^{(k)}}(x)$.
- 9 You will submit your own code for this problem; there is no auto-grader.

10 **Problem 2 (Simultaneous Localization and Mapping (SLAM) with a particle**
11 **filter, 80 points. Do not use Google Colab to do this homework).** In this prob-
12 lem, we will implement mapping and localization in an indoor environment using
13 information from an IMU and a LiDAR sensor. We have provided you data collected
14 from a humanoid named THOR that was built at Penn and UCLA
15 (<https://archive.darpa.mil/roboticschallenge/finalist/thor.html>). You can read more
16 about the hardware in this paper (<https://ieeexplore.ieee.org/document/7057369>.)



17

18 **Hardware setup of Thor** The humanoid has a Hokuyo LiDAR sensor ([https://hokuyo-](https://hokuyo-usa.com/products/lidar-obstacle-detection)
19 [usa.com/products/lidar-obstacle-detection](https://hokuyo-usa.com/products/lidar-obstacle-detection) on its head (the final version of the robot
20 had it in its chest but this is a different version); details of this are in the code (which
21 will be explained shortly). This LiDAR is a planar LiDAR sensor and returns 1080
22 readings at each instant, each reading being the distance of some physical object
23 along a ray that shoots off at an angle between $(-135, 135)$ degrees with discretiza-
24 tion of 0.25 degrees in an horizontal plane (shown as white rays in the picture).
25 We will use the position and orientation of the head of the robot to calculate the
26 orientation of the LiDAR in the body frame.

1 The second kind of observations we will use pertain to the location of the robot.
2 However, in contrast to the previous homework where we used the raw accelerometer
3 and gyroscope readings to get the orientation, we will directly use the (x, y, θ) pose
4 of the robot in the world coordinates (θ denotes yaw). These poses were created
5 presumably on the robot by running a filter on the IMU data (such estimates are
6 called odometry estimates), and just as you saw some tracking errors in the previous
7 homework, these poses will not be extremely accurate. However, we will treat them
8 conceptually the same way as we treated Vicon in the previous homework, namely
9 as a much more precise estimate of the pose of the robot that is used to check how
10 well SLAM is working.

11 **Coordinate frames** The body frame is at the top of the head (X axis pointing
12 forwards, Y axis pointing left and Z axis pointing upwards), the top of the head is at
13 a height of 1.263m from the ground. The transformation from the body frame to the
14 LiDAR frame depends upon the angle of the head (yaw) and the angle of the neck
15 (pitch) and the height of the LiDAR above the head (which is 0.15m). The world
16 coordinate frame where we want to build the map has its origin on the ground plane,
17 i.e., the origin of the body frame is at a height of 1.263m with respect to the world
18 frame at location (x, y, θ) .

19 Data and code

20 (a) **(0 points)** We have provided you 4 datasets corresponding to 4 different
21 trajectories of the robot in Towne Building at Penn. For example, dataset 0 consists
22 of two files `data/train/train_lidar0.mat` and `data/train/train_joint0.mat` which contain
23 the LiDAR readings and joint angles respectively. The functions `load_lidar_data`
24 and `load_joint_data` inside `load_data.py` read the data. You can run the function
25 `show_lidar` to see the LiDAR data. Each of the data reading functions returns
26 a data-structure where t refers to the time-stamp (in seconds) of the data, $xyth$
27 refers to (x, y, θ) pose of the LiDAR and rpy refers to Euler angles (roll, pitch,
28 yaw). The joint data contains a number of fields, but we are only interested in the
29 angle of the head and the neck at a particular time-stamp. You should read these
30 functions carefully and check the values returned by them. The dicts `joint_names`
31 and `joint_names_to_index` can be used to read off the data of a specific joint (we
32 only need the head and the neck).

33 (b) **(0 points)** Next look at the `slam.py` file provided to you. Read the code for the
34 class `map_t` and `slam_t` and the comments provided in the code very carefully. You
35 are in charge of filling in the missing pieces marked as `TODO: XXXXXX`. A sug-
36 gested order for studying this code is as follows: `slam_t.read_data`, `slam_t.init_sensor_model`,
37 `slam_t.init_particles`, `slam_t.rays2world`, `map_t.__init__`, `map_t.grid.cell_from_xy`.
38 Next, the file `utils.py` contains a few standard rigid-body transformations that
39 you will need. You should pay attention to the functions `smart_plus_2d` and
40 `smart_minus_2d` that will be used to code up the dynamics propagation step of
41 the particle filter.

1 (c) **(10 points, dynamics step)** Next look at `main.py` which has two functions
 2 `run_dynamics_step` and `run_observation_step` which act as test functions to check
 3 if the particle filter and occupancy grid update has been updated correctly. The
 4 `run_dynamics` function plots the trajectory of the robot (as given by its IMU data
 5 in the LiDAR data-structure). It also initializes 3 particles and plots all particles at
 6 different time-steps while performing the dynamics step with a very small dynamics
 7 noise; this is a very neat way of checking if dynamics propagation in the particle
 8 filter is working correctly. This function will create two plots, one for the odometry
 9 trajectory and one more for the particle trajectories, both these trajectories should
 10 match after you code up the dynamics function `slam.t.dynamics_step` correctly.

11 (d) **(20 points, observation step)** The function `run_observation_step` is used to
 12 perform the observation step of the particle filter to get an estimate of the location
 13 of the robot and updates to the occupancy grid using observations from the LiDAR.
 14 First read the comments for the function `slam.t.observation_step` carefully.

15 We first discuss the particle filtering part.

- 16 (i) Compute the head and neck position for the time t . For each particle,
 17 assuming that that particle is indeed the true position of the robot, project
 18 the LiDAR scan `slam.t.lidar[t]['scan']` into the world coordinates using the
 19 `slam.t.ray2world` function. The end points of each ray tell us which cells in
 20 the map are occupied, for each particle.
 21 (ii) In order to compute the updated weights of the particle, we need to know
 22 the likelihood of LiDAR scans given the state (our current occupancy grid
 23 in the case of SLAM). We are going to use a simple model to do so

$$\log P(\text{LiDAR scan as if the robot is at particle } p \mid m) = \sum_{ij \in O} m_{ij} \quad (1)$$

24 where O is the set of occupied cells as detected by the LiDAR scan assuming
 25 the robot is at particle p and m_{ij} is our current estimate of the binarized
 26 map (more on this below). In simple words, if the occupied cells as given
 27 by our LiDAR match the occupied cells in the binarized map created from
 28 the past observations, then we say the log-probability of particle p is large.

- 29 (iii) You will next implement the function `slam.t.update_weights` that takes
 30 the log-probability of each particle p , its previous weights, calculates the
 31 updated weights of the particles.
 32 (iv) Typically, resampling step (`slam.t.stratified_resampling`) is performed only
 33 if the effective number of particles (as computed in `slam.t.resample_particles`)
 34 falls below a certain threshold (30% in the code). Implement resampling as
 35 we discussed in the lecture notes.

36 We will now do the mapping part. We have a number of particles $p^i = (x^i, y^i, \theta^i)$
 37 that together give an estimate of the distribution of the location of the robot. For
 38 this homework, you will only use the particle with the largest weight to update
 39 the map although typically we update the map using all particles. Our goal is
 40 simple: we want to increase `map.t.log_odds` array at cells that are recorded as

1 obstacles by the LiDAR and decrease the values in all other cells. You should
2 add `slam.t.log_odds_occ` to all occupied cells and add `slam.t.log_odds_free` from
3 all cells in the map. It is also a good idea to clip the `log_odds` to like between `[-`
4 `slam.t.map.log_odds_max, slam.t.map.log_odds_max]` to prevent increasingly large
5 values in the `log_odds` array. The array `slam.t.map.cells` is a binarized version of the
6 map (which is used above to calculate the observation likelihood).

7 Check the `run_observation_step` function after you have implemented the obser-
8 vation step.

9 (e) Since the map is initialized to zero at the beginning of SLAM which results
10 in all observation log-likelihoods to be zero in (1), we need to do something special
11 for the first step. We will use the first entry in `slam.t.lidar[0]['xyth']` to get an
12 accurate pose for the robot and use its corresponding LiDAR readings to initialize
13 the occupancy grid. You can do this easily by initializing the particle filter to have
14 just one particle and simply calling the `slam.t.observation_step` as shown in `main.py`.

15 (d) **(50 points)** You will now run the full SLAM algorithm that performs one
16 dynamics step and observation step at each iteration in the function `run_slam` in
17 `main.py`. Make sure to start SLAM only after the time when you have both LiDAR
18 scans and joint readings (the two arrays start at different times). For all 4 datasets,
19 you will plot the final binarized version of the map, (x, y) location of the particle
20 in the particle filter with the largest weight at each time-step and the odometry
21 trajectory (x, y) (in a different color); this counts for 10 points each. We will use an
22 autograder for this problem where we will run your code against one test dataset
23 which will account for the other 10 points.

24 **Some Notes** This problem is much easier and shorter than it may seem. You should
25 go through these steps carefully and in the suggested order. You should make
26 sure that the results of the previous step are correct before proceeding. The two
27 functions in `main.py` to check the dynamics and observation step are very important
28 to find bugs. **Do not use Colab to do this homework.** Debugging this is much
29 easier if you use Jupyter/IPython; using `ipdb`/some other IDE (say PyCharm or
30 Spyder) is invaluable to quickly code up these kinds of problems. You do not need
31 to implement more than 100 lines of code in this problem.