

¹ Chapter 4

² Rigid-body transforms and ³ mapping

Reading

1. LaValle Chapter 3.2 for rotation matrices, Chapter 4.1-4.2 for quaternions
2. Thrun Chapter 9.1-9.2 for occupancy grids
3. OctoMap: An Efficient Probabilistic 3D Mapping Framework Based on Octrees
<http://www.arminhornung.de/Research/pub/hornung13auro.pdf>,
also see <https://octomap.github.io>.
4. Robot Operating System
<http://www.willowgarage.com/sites/default/files/icraoss09-ROS.pdf>, Optional: Lightweight Communications and Marshalling (LCM) system
<https://people.csail.mit.edu/albert/pubs/2010-huang-olson-moore-lcm-iros.pdf>
5. A Perception-Driven Autonomous Urban Vehicle
<https://april.eecs.umich.edu/media/pdfs/mitduc2009.pdf>
6. Optional reading: Thrun Chapter 10 for simultaneous localization and mapping

⁴ In the previous chapter, we looked at ways to estimate the state of the robot
⁵ in the physical world. We kept our formulation abstract, e.g., the way the robot
⁶ moves was captured by an abstract expression like $x_{k+1} = f(x_k, u_k) + \epsilon$
⁷ and observations $y_k = g(x_k) + \nu$ were similarly opaque. In order to actually
⁸ implement state estimation algorithms on real robots, we need to put concrete
⁹ functions in place of f, g .

10 This is easy to do for some robots, e.g., the robot in Problem 1 in Homework 1 moved across cells. Of course real robots are a bit more complicated,
 11 e.g., a car cannot move sideways (which is a huge headache when you parallel park). In the first half of this chapter, we will look at how to model the
 12 dynamics f using rigid-body transforms.
 13

14 The story of measurement models and sensors is similar. Although we
 15 need to write explicit formulae in place of the abstract function g . In the
 16 second half, we will study occupancy grids and dig deeper into a typical
 17 state-estimation problem in robotics, namely that of mapping the location of
 18 objects in the world around the robot.
 19

20 4.1 Rigid-Body Transformations

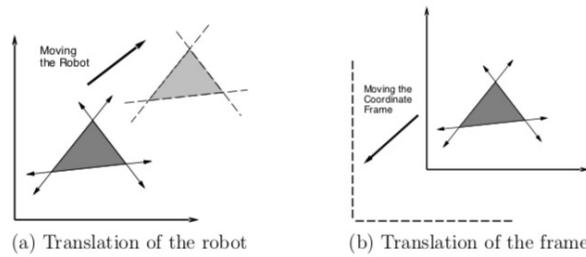
21 Let us imagine that the robot has a rigid body, we think of this as a subset
 22 $A \subset \mathbb{R}^2$. Say the robot is a disc

$$A = \{(x, y) \in \mathbb{R}^2 : x^2 + y^2 \leq 1\}.$$

23 This set A changes as the robot moves around, e.g., if the center of mass of
 24 the robot is translated by $x_t, y_t \in \mathbb{R}$ the set A changes to

$$A' = \{(x + x_t, y + y_t) : (x, y) \in A\}.$$

25 The concept of “degrees of freedom” denotes the maximum number of in-
 26 dependent parameters needed to completely characterize the transformation
 27 applied to a robot. Since the set of allowed values (x_t, y_t) is a two-dimensional
 28 subset of \mathbb{R}^2 , then the degrees of freedom available to a translating robot is
 29 two.



30 As the above figure shows, there are two ways of thinking about this transfor-
 31 mation. We can either think of the robot transforming while the co-ordinate
 32 frame of the world is fixed, or we can think of it as the robot remaining sta-
 33 tionary and the co-ordinate frame undergoing a translation. The second style
 34 is useful if you want to imagine things from the robot’s perspective. But the
 35 first one feels much more natural and we will therefore exclusively use the
 36 first notion.
 37

38 If the same robot if it were rotated counterclockwise by some angle
 39 $\theta \in [0, 2\pi]$, we would map

$$(x, y) \mapsto (x \cos \theta - y \sin \theta, x \sin \theta + y \cos \theta).$$

⁴⁰ Such a map can be written as multiplication by a 2×2 rotation matrix

$$R(\theta) = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix}. \quad (4.1)$$

⁴¹ to get

$$\begin{bmatrix} x \cos \theta - y \sin \theta \\ x \sin \theta + y \cos \theta \end{bmatrix} = R(\theta) \begin{bmatrix} x \\ y \end{bmatrix}.$$

⁴² The transformed robot is thus given by

$$A' = \left\{ R \begin{bmatrix} x \\ y \end{bmatrix} : (x, y) \in A \right\}.$$

⁴³ If we perform both rotation and translation, we can the transformation using a
⁴⁴ single matrix

$$T = \begin{bmatrix} \cos \theta & -\sin \theta & x_t \\ \sin \theta & \cos \theta & y_t \\ 0 & 0 & 1 \end{bmatrix} \quad (4.2)$$

⁴⁵ and this transformation looks like

$$\begin{bmatrix} x \cos \theta - y \sin \theta + x_t \\ x \sin \theta + y \cos \theta + y_t \\ 1 \end{bmatrix} = T \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}.$$

⁴⁶ The point $(x, y, 1) \in \mathbb{R}^3$ is called homogeneous coordinate space correspond-
⁴⁷ ing to $(x, y) \in \mathbb{R}^2$ and the matrix T is called a *homogeneous transformation*
⁴⁸ *matrix*. The peculiar names comes from the fact that even if the matrix T
⁴⁹ maps rotations and translations of rigid bodies $A \subset \mathbb{R}^2$, it is just a linear
⁵⁰ transformation of the point $(x, y, 1)$ if viewed in the larger space \mathbb{R}^3 .

⁵¹ **Rigid-body transformations** The transformations $R \in \mathbb{R}^{2 \times 2}$ or $T \in \mathbb{R}^{3 \times 3}$
⁵² are called rigid-body transformations. Mathematically, it means that they do
⁵³ not cause the distance between any two points inside the set A to change. Rigid-
⁵⁴ body transformations are what are called an orthogonal group in mathematics.

⁵⁵ **A group** is a mathematical object which imposes certain conditions upon
⁵⁶ how two operations, e.g., rotations, can be composed together. For instance,
⁵⁷ if G is the group of rotations, then (i) the composition of two rotations is a
⁵⁸ rotation, we say that it satisfies closure $R(\theta_1)R(\theta_2) \in G$, (ii) rotations are
⁵⁹ associative

$$R(\theta_1) \{R(\theta_2)R(\theta_3)\} = \{R(\theta_1)R(\theta_2)\} R(\theta_3),$$

⁶⁰ and, (iii) there exists an identity and inverse rotation

$$R(0), R(-\theta) \in G.$$

⁶¹ **An orthogonal group** is a group whose operations preserve distances in
⁶² Euclidean space, i.e., $g \in G$ is an element of the group that acts on two points

❶ It is important to remember that T represents rotation *followed* by a translation, not the other way around.

63 $x, y \in \mathbb{R}^d$ then

$$\|g(x) - g(y)\| = \|x - y\|.$$

64 If we identify the basis in Euclidean space to be the set of orthonormal vec-
65 tors $\{e_1, \dots, e_d\}$, then equivalently, the orthogonal group $O(d)$ is the set of
66 orthogonal matrices

$$O(d) := \{O \in \mathbb{R}^{d \times d} : OO^\top = O^\top O = I\}.$$

67 This implies that the square of the determinant of any element $a \in O(d)$ is 1,
68 i.e., $\det(a) = \pm 1$.

❶ Check that any rotation matrix R belongs to an orthogonal group.

69 **The Special Orthogonal Group** is a sub-group of the orthogonal group
70 where the determinant of each element is +1. You can see that rotations are a
71 special orthogonal group. We denote rotations of objects in \mathbb{R}^2 as

$$SO(2) := \{R \in \mathbb{R}^{2 \times 2} : R^\top R = RR^\top = I, \det(R) = 1\}. \quad (4.3)$$

72 Each group element $g \in SO(2)$ denotes a rotation of the XY -plane about the
73 Z -axis. The group of 3D rotations is called the Special Orthogonal Group
74 $SO(3)$ and is defined similarly

$$SO(3) := \{R \in \mathbb{R}^{3 \times 3} : R^\top R = RR^\top = I, \det(R) = 1\}. \quad (4.4)$$

75 **The Special Euclidean Group $SE(2)$** is simply a composition of a 2D
76 rotation $R \in SO(2)$ and a 2D translation $\mathbb{R}^2 \ni v \equiv (x_t, y_t)$

$$SE(2) = \left\{ \begin{bmatrix} R & v \\ 0 & 1 \end{bmatrix} : R \in SO(2), v \in \mathbb{R}^2 \right\} \subset \mathbb{R}^{3 \times 3}. \quad (4.5)$$

77 The Special Euclidean Group $SE(3)$ is defined similarly as

$$SE(3) = \left\{ \begin{bmatrix} R & v \\ 0 & 1 \end{bmatrix} : R \in SO(3), v \in \mathbb{R}^3 \right\} \subset \mathbb{R}^{4 \times 4}; \quad (4.6)$$

78 again, remember that it is rotation *followed* by a translation.

79 4.1.1 3D transformations

80 Translations and rotations in 3D are conceptually similar to the two-dimensional
81 case; however the details appear a bit more difficult because rotations in 3D
82 are more complicated.

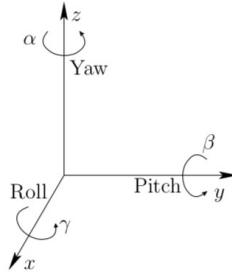


Figure 4.1: Any three-dimensional rotation can be described as a sequence of rotations about each of the cardinal axes. We usually give these specific names: rotation about the Z -axis is called yaw, rotation about the X -axis is called roll and rotation about Y -axis is called pitch. You should commit this picture and these names to memory because it will be of enormous to think about these rotations intuitively.

⁸³ **Euler angles** We know that a pure counter-clockwise rotation about one of
⁸⁴ the axes is written in terms of a matrix, say yaw of α -radians about the Z -axis

$$R_z(\alpha) = \begin{bmatrix} \cos \alpha & -\sin \alpha & 0 \\ \sin \alpha & \cos \alpha & 0 \\ 0 & 0 & 1 \end{bmatrix}.$$

⁸⁵ Notice that this is a 3×3 matrix that keeps the Z -coordinate unchanged and
⁸⁶ only affects the other two coordinates. Similarly we have for pitch (β about
⁸⁷ the Y -axis) and roll (γ about the X -axis)

$$R_y(\beta) = \begin{bmatrix} \cos \beta & 0 & \sin \beta \\ 0 & 1 & 0 \\ -\sin \beta & 0 & \cos \beta \end{bmatrix}, R_x(\gamma) = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos \gamma & -\sin \gamma \\ 0 & \sin \gamma & \cos \gamma \end{bmatrix}.$$

⁸⁸ By convention, a rotation matrix in three dimensions is understood as a se-
⁸⁹ quential application of rotations, *first roll, then pitch, and then yaw*

$$\mathbb{R}^{3 \times 3} = R(\alpha, \beta, \gamma) = R_z(\alpha)R_y(\beta)R_x(\gamma). \quad (4.7)$$

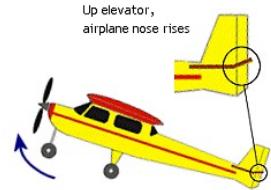
⁹⁰ The angles (α, γ, β) are called Euler angles. Imagine how the body frame of
⁹¹ the robot changes as successive rotations are applied. If you were sitting in a
⁹² car, a pure yaw would be similar to the car turning left; the Z -axis correspond-
⁹³ ing to this yaw would however only be pointing straight up perpendicular to
⁹⁴ the ground if you had not performed a roll/pitch before. If you had, the Z -axis
⁹⁵ of the body frame with respect to the world will be tilted.

⁹⁶ Another important thing to note is that one single parameter determines
⁹⁷ all possible rotations about one axis, i.e., $SO(2)$. But three Euler angles are
⁹⁸ used to parameterize general rotations in three-dimensions. You can watch
⁹⁹ <https://www.youtube.com/watch?v=3Zjf95Jw2UE> to get more intuition about
¹⁰⁰ Euler angles.

❶ Here is how I remember these names. Say you are driving a car, usually in robotics we take the X -axis to be longitudinally forward, the Y -axis is your left hand if you are in the driver's seat and the Z -axis points up by the right-hand thumb rule. Roll



is what a dog does when it rolls, it rotates about the X -axis. Pitch is what a plane



does when it takes off, its nose lifts up and it rotates about the Y -axis. Yaw is the one leftover.

101 **Rotation matrices to Euler angles** We can back-calculate the Euler angles
102 from a rotation matrix as follows. Given an arbitrary matrix

$$R = \begin{bmatrix} r_{11} & r_{12} & r_{13} \\ r_{21} & r_{22} & r_{23} \\ r_{31} & r_{32} & r_{33} \end{bmatrix},$$

103 we set

$$\begin{aligned} \alpha &= \tan^{-1}(r_{21}/r_{11}) \\ \beta &= \tan^{-1}\left(-r_{31}/\sqrt{r_{32}^2 + r_{33}^2}\right) \\ \gamma &= \tan^{-1}(r_{32}/r_{33}). \end{aligned} \quad (4.8)$$

104 For each angle, the corresponding quadrant for the Euler angle is determined
105 using the signs of the numerator and the denominator. So you should use
106 the function atan2 in Python/C++ to implement these expressions correctly.
107 Notice that some of the expressions have r_{11} and r_{33} in the denominator, this
108 means that we need $r_{11} = \cos \alpha \cos \beta = 0$ and $r_{33} = \cos \beta \cos \gamma = 0$. Euler
109 angles do not completely parametrize the group of rotations $SO(3)$. This is
110 really an indicator that this particular physical rotation can be parameterized
111 in two different ways using Euler angles, so the map from rotation matrices to
112 Euler angles is not unique.

113 **Homogeneous coordinates in three dimensions** Just like the 2D case, we
114 can define a 4×4 matrix that transforms points $(x, y, z) \in \mathbb{R}^3$ to their new
115 locations after a rotation by Euler angles (α, β, γ) and a translation by a vector
116 $v = (x_t, y_t, z_t) \in \mathbb{R}^3$

$$T = \begin{bmatrix} R(\alpha, \beta, \gamma) & v \\ 0 & 1 \end{bmatrix}.$$

117 4.1.2 Rodrigues' formula: an alternate view of rotations

118 Consider a point $r(t) \in \mathbb{R}^3$ that is being rotated about an axis denoted by a
119 unit vector $\omega \in \mathbb{R}^3$ with an angular velocity of 1 radian/sec. The instantaneous
120 linear velocity of the head of the vector is

$$\dot{r}(t) = \omega \times r(t) \equiv \hat{\omega}r(t) \quad (4.9)$$

121 where the \times denotes the cross-product of the two vectors $a, b \in \mathbb{R}^3$

$$a \times b = \begin{bmatrix} a_2 b_3 - a_3 b_2 \\ a_3 b_1 - a_1 b_3 \\ a_1 b_2 - a_2 b_1 \end{bmatrix}$$

122 which we can equivalently denote as a matrix vector multiplication $a \times b = \hat{a}b$
123 where

$$\hat{a} = \begin{bmatrix} 0 & -a_3 & a_2 \\ a_3 & 0 & -a_1 \\ -a_2 & a_1 & 0 \end{bmatrix}. \quad (4.10)$$

¹²⁴ is a skew-symmetric matrix. The solution of the differential equation (4.9) at
¹²⁵ time $t = \theta$ is

$$r(\theta) = \exp(\hat{\omega}\theta) r(0)$$

¹²⁶ where the matrix exponential of a matrix A is defined as

$$\exp(A) = I + A + \frac{A^2}{2!} + \frac{A^3}{3!} + \dots = \sum_{k=0}^{\infty} \frac{A^k}{k!}.$$

¹²⁷ This is an interesting observation: a rotation about a fixed axis ω by an angle θ
¹²⁸ can be represented by the matrix

$$R = \exp(\hat{\omega}\theta).$$

¹²⁹ You can check that this matrix is indeed a rotation by showing that $R^\top R = I$
¹³⁰ and that $\det(R) = +1$. We can expand the matrix exponential and collect odd
¹³¹ and even powers of $\hat{\omega}$ to get

$$R = I + \sin \theta \hat{\omega} + (1 - \cos \theta) \hat{\omega}^2. \quad (4.11)$$

¹³² which is the Rodrigues' formula that relates the angle θ and the axis ω to the
¹³³ rotation matrix. We can also go in the opposite direction, i.e., given a matrix
¹³⁴ R calculate what angle θ and axis ω it corresponds to using

$$\begin{aligned} \cos \theta &= \frac{\text{tr}(R) - 1}{2} \\ \hat{\omega} &= \frac{R - R^\top}{2 \sin \theta}. \end{aligned} \quad (4.12)$$

¹³⁵ 4.2 Quaternions

¹³⁶ We know two ways to think about rotations: we can either think in terms of the
¹³⁷ three Euler angles (α, β, γ) , or we can consider a rotation matrix $R \in \mathbb{R}^{3 \times 3}$.
¹³⁸ We also know ways to go to and fro between these two forms with the caveat
¹³⁹ that solving for Euler angles using (4.8) may be degenerate in some cases.
¹⁴⁰ While rotation matrices are the most general representation of rotations, using
¹⁴¹ them in computer code is cumbersome (it is, after all, a matrix of 9 elements),
¹⁴² imagine an EKF where the state is a rotation matrix. You can certainly
¹⁴³ implement the same filter using Euler angles but doing so requires special
¹⁴⁴ care due to the degeneracies. Quaternions are a neat way to avoid both these
¹⁴⁵ problems, they parametrize the space of rotations using 4 numbers and do not
¹⁴⁶ have degeneracies.

¹⁴⁷ The central idea behind quaternions is Euler's theorem which says that
¹⁴⁸ any 3D rotation can be considered as a pure rotation by an angle $\theta \in \mathbb{R}$ about
¹⁴⁹ an axis given by the unit vector ω . This is the result that we also exploited in
¹⁵⁰ Rodrigues' formula.

❶ Quaternions were invented by British mathematician William Rowan Hamilton while walking along a bridge with his wife. He was quite excited by this discovery and promptly graffitied the expression into the stone of the bridge

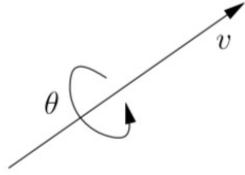


Figure 4.2: Any rotation in 3D can be represented using a unit vector ω and an angle $\theta \in \mathbb{R}$. Notice that there are two ways to encode the same rotation, the unit vector $-\omega$ and angle $2\pi - \theta$ would give the same rotation. Mathematicians say this as quaternions being a double-cover of $SO(3)$.

151 A quaternion q as a four-dimensional vector $q \equiv (u_0, u_1, u_2, u_3)$ and we
152 write it as

$$\begin{aligned} q &\equiv (u_0, u), \text{ or} \\ q &= u_0 + u_1 i + u_2 j + u_3 k, \end{aligned} \tag{4.13}$$

153 with i, j, k being three “imaginary” components of the quaternion with “complex-
154 numbers like” relationships

$$i^2 = j^2 = k^2 = ijk = -1. \tag{4.14}$$

155 It follows from these relationships that

$$ij = -ji = k, ki = -ik = j, \text{ and } jk = -jk = i.$$

156 Although you may be tempted to think about this, these imaginary components
157 i, j, k have no relationship with the square roots of negative unity used to
158 define standard complex numbers. You should simply think of the quaternion
159 as a four-dimensional vector in some space. A unit quaternion, i.e., one with

$$u_0^2 + u_1^2 + u_2^2 + u_3^2 = 1,$$

160 is special: unit quaternions can be used to represent rotations in 3D.

161 **Quaternion to axis-angle representation** The quaternion $q = (u_0, u)$ cor-
162 responds to a counterclockwise rotation of angle θ about a unit vector ω
163 where θ and ω are such that

$$u_0 = \cos \frac{\theta}{2}, \text{ and } u = \sin \frac{\theta}{2} \omega. \tag{4.15}$$

164 So given an axis-angle representation of rotation like in Rodrigues' formula
165 (θ, ω) we can write the quaternion as

$$q = \left(\cos \frac{\theta}{2}, \sin \frac{\theta}{2} \omega \right).$$

¹⁶⁶ Using this, we can also compute the inverse of a quaternion (rotation of angle
¹⁶⁷ θ about the opposite axis $-\omega$) as

$$q^{-1} := \left(\cos \frac{\theta}{2}, -\sin \frac{\theta}{2} \omega \right).$$

¹⁶⁸ The inverse quaternion is therefore the quaternion where all entries except the
¹⁶⁹ first have their signs flipped.

¹⁷⁰ **Multiplication of quaternions** Just like two rotation matrices multiply together to give a new rotation (rotations are a group), since quaternions are
¹⁷¹ also a representation for the group of rotations, we can also multiply two
¹⁷² quaternions $q_1 = (u_0, u)$, $q_2 = (v_0, v)$ together using the quaternion identities
¹⁷³ for i, j, k in (4.14) to get a new quaternion

$$q_1 q_2 \equiv (u_0, u) \cdot (v_0, v) = (u_0 v_0 - u^\top v, u_0 v + v_0 u + u \times v).$$

¹⁷⁵

¹⁷⁶ **Pure quaternions** A pure quaternion is a quaternion with a zero scalar value
¹⁷⁷ $u_0 = 0$. This is very useful to simply store a standard 3D vector $u \in \mathbb{R}^3$ as a
¹⁷⁸ quaternion $(0, u)$. We can then rotate points easily between different frames
¹⁷⁹ as follows. Given a vector $x \in \mathbb{R}^3$ we can form a quaternion $(0, x)$ and show
¹⁸⁰ that

$$q \cdot (0, x) \cdot q^* = (0, R(q)x). \quad (4.16)$$

¹⁸¹ where $q^* = (u_0, -u)$ is the conjugate quaternion of $q = (u_0, u)$; the conjugate
¹⁸² is the same as the inverse for unit quaternions. Notice how the right-hand side
¹⁸³ is the vector $R(q)x$ corresponding to the vector x rotation by a matrix $R(q)$.

¹⁸⁴ **Quaternions to rotation matrix** The rotation matrix corresponding to a
¹⁸⁵ quaternion is

$$\begin{aligned} R(q) &= (u_0^2 - u^\top u) I_{3 \times 3} + 2 \frac{u_0 u}{\|u\|} + 2 u u^\top \\ &= \begin{bmatrix} 2(u_0^2 + u_1^2) - 1 & 2(u_1 u_2 - u_0 u_3) & 2(u_1 u_3 + u_0 u_2) \\ 2(u_1 u_2 + u_0 u_3) & 2(u_0^2 + u_2^2) - 1 & 2(u_2 u_3 - u_0 u_1) \\ 2(u_1 u_3 - u_0 u_2) & 2(u_2 u_3 + u_0 u_1) & 2(u_0^2 + u_3^2) - 1 \end{bmatrix}. \end{aligned} \quad (4.17)$$

¹⁸⁶ Using this you can show the identity that rotation matrix corresponding to the
¹⁸⁷ product of two quaternions is the product of the individual rotation matrices

$$R(q_1 q_2) = R(q_1) R(q_2).$$

❶ Quaternions belong to a larger group than rotations called the Symplectic Group $Sp(1)$.

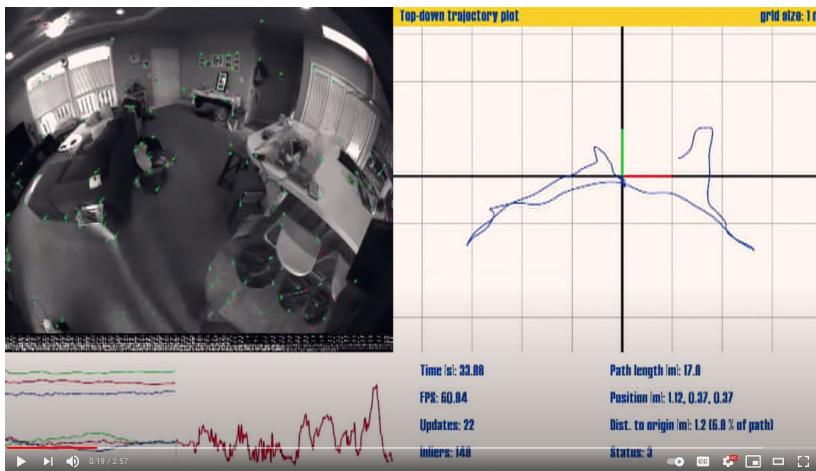
¹⁸⁸ **Rotation matrix to quaternion** We can also go in the reverse direction.
¹⁸⁹ Given a rotation matrix R , the quaternion is

$$\begin{aligned}
 u_0 &= \frac{1}{2} \sqrt{r_{11} + r_{22} + r_{33} + 1} \\
 \text{if } u_0 \neq 0, \quad u_1 &= \frac{r_{32} - r_{23}}{4u_0} \\
 u_2 &= \frac{r_{13} - r_{31}}{4u_0} \\
 u_3 &= \frac{r_{21} - r_{12}}{4u_0} \\
 \text{if } u_0 = 0, \quad u_1 &= \frac{r_{13}r_{12}}{\sqrt{r_{12}^2r_{13}^2 + r_{12}^2r_{23}^2 + r_{13}^2r_{23}^2}} \\
 u_2 &= \frac{r_{12}r_{23}}{\sqrt{r_{12}^2r_{13}^2 + r_{12}^2r_{23}^2 + r_{13}^2r_{23}^2}} \\
 u_3 &= \frac{r_{13}r_{23}}{\sqrt{r_{12}^2r_{13}^2 + r_{12}^2r_{23}^2 + r_{13}^2r_{23}^2}}.
 \end{aligned} \tag{4.18}$$

¹⁹⁰ 4.3 Occupancy Grids

¹⁹¹ Rotation matrices and quaternions let us capture the dynamics of a rigid robot
¹⁹² body. We will next look at how to better understand observations.

¹⁹³ **What is location and what is mapping?** Imagine a robot that is moving
¹⁹⁴ around in a house. A natural representation of the state of this robot is the 3D lo-
¹⁹⁵ cation of all the interesting objects in the room, e.g., <https://www.youtube.com/watch?v=Qe10ExwzCqk>.
¹⁹⁶ At each time-instant, we record an observation from our sensor (in this case,
¹⁹⁷ a camera) that indicates how far an object is from the robot. This helps us
¹⁹⁸ discover the location of the objects in the room. After gathering enough ob-
¹⁹⁹ servations, we would have created a *map* of the entire house. This map is the
²⁰⁰ set of positions of all interesting objects in the room. Such a map is called a
²⁰¹ “feature map”, these are all the green points in the image below



²⁰³ The main point to understand about feature map is that we can hand over
²⁰⁴ this map to another robot that comes to the same house. The robot compares

205 images from its camera and if it finds one of the objects inside the map, it
 206 can get an estimate of its location/orientation in the room with respect to the
 207 known location of the object in the map. The map is just a set of “features”
 208 that help identify salient objects in the room (objects which can be easily
 209 detected in images and relatively uniquely determine the location inside the
 210 room). The second robot using this map to estimate its position/orientation in
 211 the room is called the *localization problem*. We already know how to solve
 212 the localization problem using filtering.

213 The first robot was solving a harder problem called *Simultaneous Local-
 214 ization And Mapping (SLAM)*: namely that of discovering the location of both
 215 itself and the objects in the house. This is a very important and challenging
 216 problem in robots but we will not discuss it further. MEAM 620 digs deeper
 217 into it.

In this section, we will solve a part of the SLAM problem, namely the mapping problem. We will assume that we know the position/orientation of the robot in the 3D world, and want to build a map of the objects in the world. We will discuss grid maps, which are a more crude way of representing maps than feature maps but can be used easily even if there are lots of objects.

218 **Grid maps** We will first discuss two-dimensional grid maps, they look as
 219 follows.

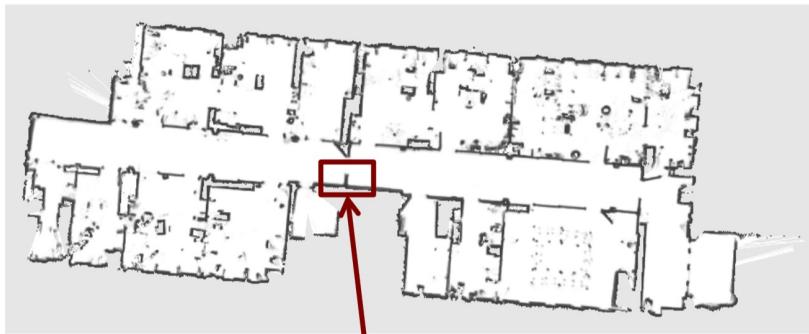
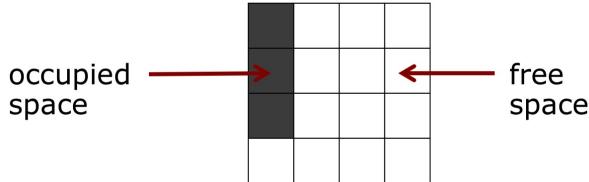


Figure 4.3: A grid map (also called an occupancy grid) is a large gray-scale image, each pixel represents a cell in the physical world. In this picture, cells that are occupied are colored black and empty cells represent free space. A grid map is a useful representation for a robot to localize in this house using observations from its sensors and comparing those to the map.

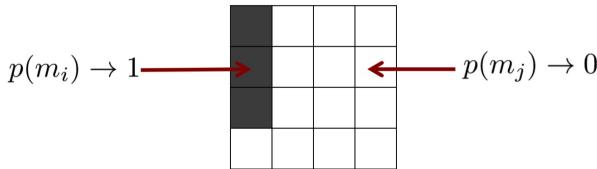
220 To get a quick idea of what we want to do, you can watch the mapping
 221 being performed in <https://www.youtube.com/watch?v=JJhEkIA1xSE>. We are
 222 interested in learning such maps from the observations that a robot collects as
 223 it moves around the physical space. Let us make two simplifying assumptions.

224 **Assumption 1: each cell is either free or occupied**



225

226 This is neat: we can now model each cell as a binary random variable that
227 indicates occupancy. Let the probability that the cell m_i be occupied be $p(m_i)$

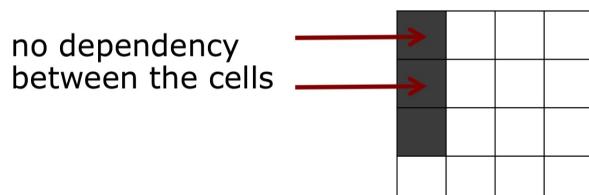


228

229 If we have $p(m_i) = 0$, then the cell is not occupied and if we have $p(m_i) = 1$,
230 then the cell is occupied. A priori, we do not know the state of the cell so we
231 will set the prior probability to be $p(m_i) = 0.5$.

232 **Assumption 2: the world is static** Objects in the world do not move. This
233 is reasonable if we are interested in estimating in building a map of the walls
234 inside the room. Note that it is not a reasonable assumption if there are moving
235 people inside the room. We will see a clever hack where the Bayes rule helps
236 automatically disregard such moving objects in this section.

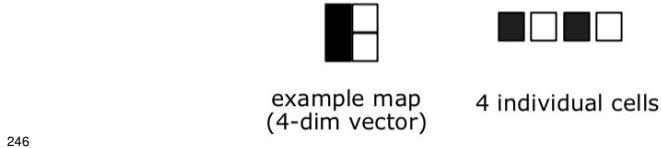
237 **Assumption 3: cells are independent of each other** This is another drastic
238 simplification. The state of our system is the occupancy of each cell in the
239 grid map. We assume that before receiving any observations, the occupancy of
240 each individual cell is independent; it is a Bernoulli variable with probability
241 1/2 since we have assumed the prior to be uniform in Assumption 1.



242

243 This means that if cells in the map are denoted by a vector $m = (m_1, \dots,)$,
244 then the probability of the cells being occupied/not-occupied can be written as

$$245 \quad p(m) = \prod_i p(m_i). \quad (4.19)$$



246

247 4.3.1 Estimating the map from the data

248 Say that the robot pose (position and orientation) is given by the sequence
 249 x_1, \dots, x_k . While proceeding along this sequence, the robot receives obser-
 250 vations y_1, \dots, y_k . Our goal is to estimate the state of each cell $m_i \in \{0, 1\}$
 251 (aka “the map” $m = (m_1, m_2, \dots,)$)

$$P(m | x_1, \dots, x_k, y_1, \dots, y_k) = \prod_i P(m_i | x_1, \dots, x_k, y_1, \dots, y_k). \quad (4.20)$$

252 This is called the “static state” Bayes filter and is conceptually exactly the
 253 same as the recursive application of Bayes rule in Chapter 2 for detecting
 254 whether the door was open or closed.

255 We will use a short form to keep the notation clear

$$y_{1:k} = (y_1, y_2, \dots, y_k);$$

256 the quantity $x_{1:k}$ is defined similarly. As usual we will use a recursive Bayes
 257 filter to compute this probability as follows.

$$\begin{aligned} P(m_i | x_{1:k}, y_{1:k}) &\stackrel{\text{Bayes rule}}{=} \frac{P(y_k | m_i, y_{1:k-1}, x_{1:k}) P(m_i | y_{1:k-1}, x_{1:k})}{P(y_k | y_{1:k-1}, x_{1:k})} \\ &\stackrel{\text{Markov}}{=} \frac{P(y_k | m_i, x_k) P(m_i | y_{1:k-1}, \cancel{x_{1:k-1}})}{P(y_k | y_{1:k-1}, x_{1:k})} \\ &\stackrel{\text{Bayes rule}}{=} \frac{P(m_i | y_k, x_k) P(y_k | x_k) P(m_i | y_{1:k-1}, x_{1:k-1})}{P(m_i | x_k) P(y_k | y_{1:k-1}, x_{1:k})} \\ &\stackrel{\text{Markov}}{=} \frac{P(m_i | y_k, x_k) P(y_k | x_k) P(m_i | y_{1:k-1}, x_{1:k-1})}{P(m_i) P(y_k | y_{1:k-1}, x_{1:k})}. \end{aligned} \quad (4.21)$$

258 We have a similar expression for the opposite probability

$$P(\neg m_i | x_{1:k}, y_{1:k}) = \frac{P(\neg m_i | y_k, x_k) P(y_k | x_k) P(\neg m_i | y_{1:k-1}, x_{1:k-1})}{P(\neg m_i) P(y_k | y_{1:k-1}, x_{1:k})}.$$

259 Let us take the ratio of the two to get

$$\begin{aligned} \frac{P(m_i | x_{1:k}, y_{1:k})}{P(\neg m_i | x_{1:k}, y_{1:k})} &= \underbrace{\frac{P(m_i | y_k, x_k)}{P(\neg m_i | y_k, x_k)} \frac{P(m_i | y_{1:k-1}, x_{1:k-1})}{P(\neg m_i | y_{1:k-1}, x_{1:k-1})}}_{\text{uses observation } y_k} \underbrace{\frac{P(\neg m_i)}{P(m_i)}}_{\text{prior}} \\ &= \underbrace{\frac{P(m_i | y_k, x_k)}{1 - P(m_i | y_k, x_k)}}_{\text{uses observation } y_k} \underbrace{\frac{P(m_i | y_{1:k-1}, x_{1:k-1})}{1 - P(m_i | y_{1:k-1}, x_{1:k-1})}}_{\text{recursive term}} \underbrace{\frac{1 - P(m_i)}{P(m_i)}}_{\text{prior}}. \end{aligned} \quad (4.22)$$

260 This is called the odds ratio. Notice that the first term uses the latest obser-
 261 vation y_k , the second term can be updated recursively because it is a similar

262 expression as the left-hand side and the third term is a prior probability of the
 263 cell being occupied/non-occupied. Let us rewrite this formula using the log-
 264 odds-ratio that makes implementing it particularly easy. The log-odds-ratio of
 265 the probability $p(x)$ of a binary variable x is defined as

$$l(x) = \log \frac{p(x)}{1 - p(x)}, \text{ and } p(x) = 1 - \frac{1}{1 + e^{l(x)}}.$$

266 The product in (4.22) now turns into a sum as

$$l(m_i | y_{1:k}, x_{1:k}) = l(m_i | y_k, x_k) + l(m_i | y_{1:k-1}, x_{1:k-1}) - l(m_i). \quad (4.23)$$

267 This expression is used to update the occupancy of each cell. The term

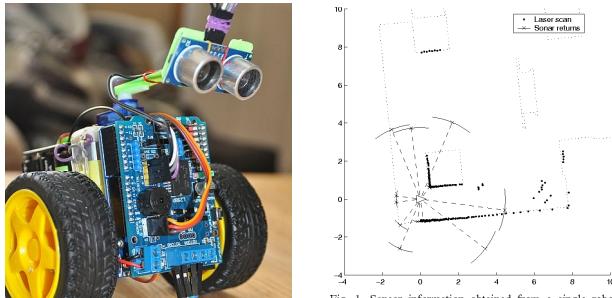
$$\text{sensor model} = l(m_i | y_k, x_k)$$

268 is different for different sensors and we will investigate it next.

💡 We assumed that the map was static. Can you think of why (4.23) automatically lets us handle some moving objects? Think of what the prior odds $l(m_i)$ does to the log-odds-ratio $l(m_i | y_{1:k}, x_{1:k})$.

269 4.3.2 Sensor models

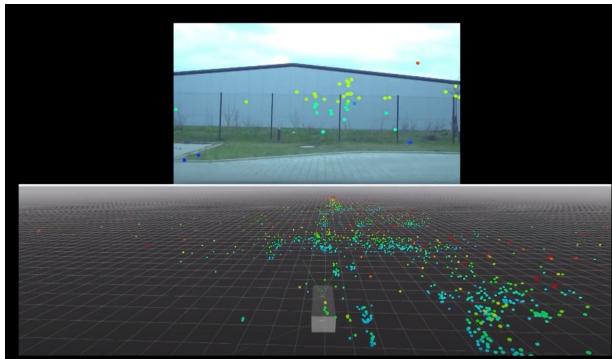
270 **Sonar** This works by sending out an ultrasonic chirp and measuring the time
 271 between emission and reception of the signal. The time gives an estimate of
 272 the distance of an object to the robot.



273 The figure above shows a typical sonar sensor (the two “eyes”) on a low-cost
 274 robot. Data from the sensor is shown on the right, a sonar is a very low
 275 resolution sensor and has a wide field of view, say 15 degrees, i.e., it cannot
 276 differentiate between objects that are within 15 degrees of each other and
 277 registers them as the same point. Sophisticated sonar technology is used today
 278 in marine environments (submarines, fish finders, detecting mines etc.).

280 **Radar** works in much the same way as a sonar except that it uses pulses
 281 of radio waves and measures the phase difference between the transmitted
 282 and the received signal. This is a very versatile sensor (it was invented by the
 283 US army to track planes and missiles during World War II) but is typically
 284 noisy and requires sophisticated processing to be used for mainstream robotics.
 285 Autonomous cars, collision warning systems on human-driven cars, weather
 286 sensing, and certainly the military use the radar today. The following picture
 287 and the video https://www.youtube.com/watch?v=hwKUcu_7F9E will give
 288 you an appreciation of the kind of data that a radar records. Radar is a very

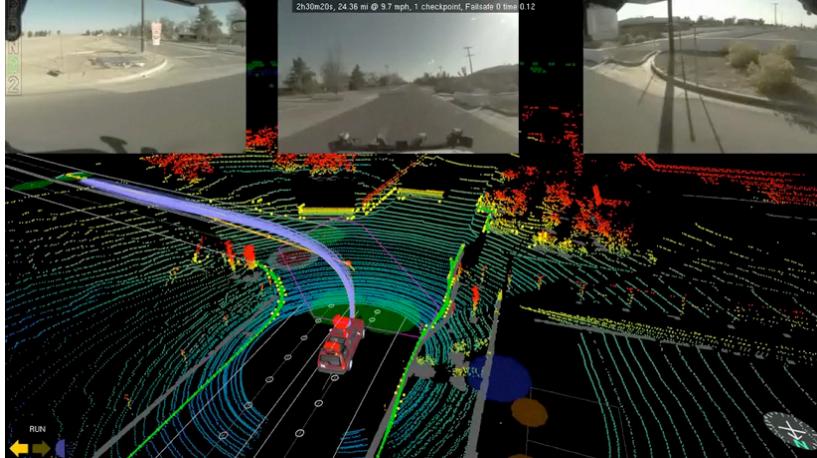
289 long range sensor (typically 150 m) and works primarily to detect metallic
 290 objects.



291

LiDAR LiDAR, which is short for Light Detection and Ranging, (<https://en.wikipedia.org/wiki/Lidar>) is a portmanteau of light and radar. It is a sensor that uses a pulsed laser as the source of illumination and records the time it takes (nanoseconds typically) for the signal to return to the source. See <https://www.youtube.com/watch?v=NZKvf1cXe8s> for how the data from a typical LiDAR (Velodyne) looks like. While a Velodyne contains an intricate system of rotating mirrors and circuitry to measure time elapsed, there are new solid state LiDARs that are rapidly evolving to match the needs of the autonomous driving industry. Most LiDARs have a usable range of about 100 m.

300



301

A typical autonomous car This is a picture of MIT's entry named Talos to the DARPA Urban Challenge ([https://en.wikipedia.org/wiki/DARPA_Grand_Challenge_\(2007\)](https://en.wikipedia.org/wiki/DARPA_Grand_Challenge_(2007))) which was a competition where teams had to traverse a 60 mile urban route within 6 hours, while obeying traffic laws, understanding incoming vehicles etc. Successful demonstrations by multiple teams led by (CMU, Stanford, Odin College, MIT, Penn and Cornell) in this competition jump-started the wave of autonomous driving. While the number of sensors necessary to drive well has come down (Tesla famously does not like to use LiDARs and rely exclusively on cameras and radars), the type of sensors and challenges associated with them remain essentially the same.

❶ Waymo's autonomous car



312



313 4.3.3 Back to sensor modeling

314 Let us go back to understanding our sensor model $l(m_i | y_k, x_k)$ where m_i
 315 is a particular cell of the occupancy grid, y_k and x_k are the observations and
 316 robot position/orientation at time k .

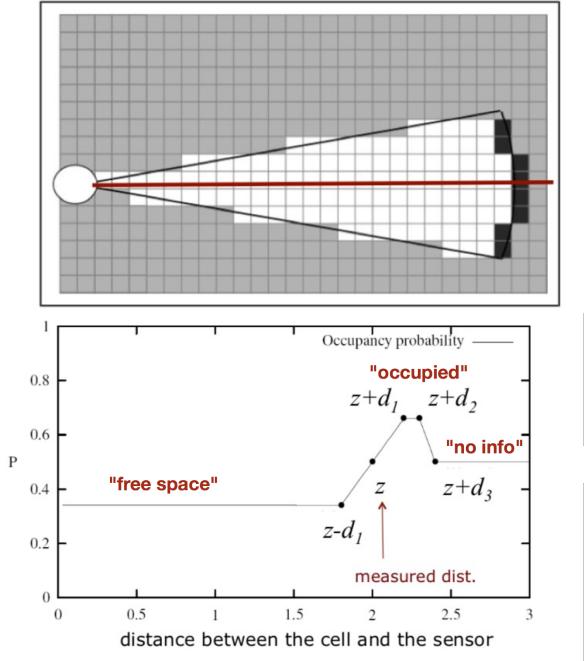
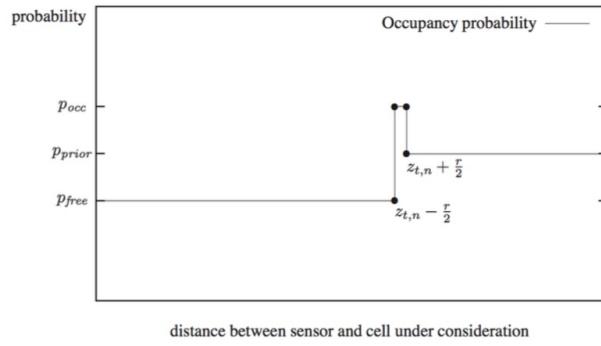


Figure 4.4: **Model for sonar data.** (Top) A sonar gives one real-valued reading corresponding to the distance measured along the red axis. (Bottom) if we travel along the optical axis, the occupancy probability $P(m_i | y_k = z, x_k)$ can be modeled as a spike around the measured value z . It is very important to remember that range sensors such as sonar gives us three kinds of information about this ray: (i) all parts of the environment up to $\approx z$ are *unoccupied* (otherwise we would not have recorded z), (ii) there is some object at z which resulted in the return, (iii) but we do not know anything about what is behind z . So incorporating a measurement y_k from a sonar/radar/lidar involves not just updating the cell which corresponds to the return, but also updating the occupancy probabilities of every grid cell along the axis.



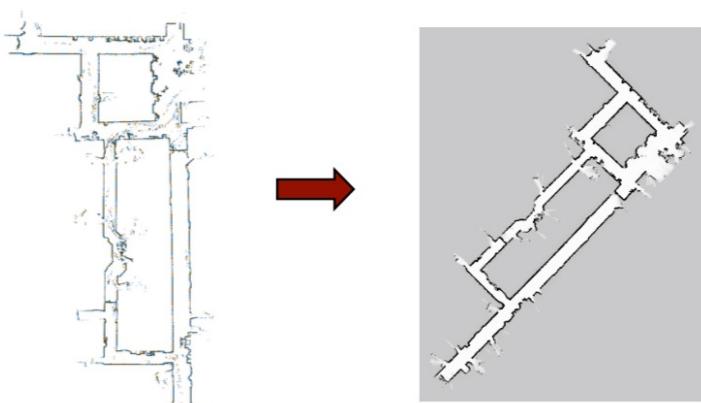
Figure 4.5: (Left) A typical occupancy grid created using a sonar sensor by updating the log-odds-ratio $l(m_i | x_{1:k}, y_{1:k})$ for all cells i for multiple time-steps k . At the end of the map building process, if $l(m_i | x_{1:k}, y_{1:k}) > 0$ for a particular cell, we set its occupancy to 1 and to zero otherwise, to get the maximum-likelihood estimate of the occupancy grid on the right.

317 **LiDAR model** When we say that a LiDAR is a more accurate sensor than
318 the sonar, what we really mean is that the sensor model $P(m_i | y_k, x_k)$ looks
319 as follows.



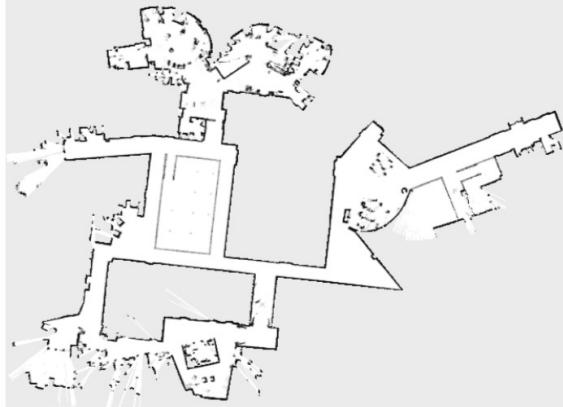
320

321 As a result, we can create high-resolution occupancy grids using a LiDAR.



322

Example: MIT CSAIL 3rd Floor



323

❸ How will you solve the localization problem given the map? In other words, if we know the occupancy grid of a building as estimated in a prior run, and we now want to find the position/orientation of the robot traveling in this building, how show we use these sensors?

324 4.4 3D occupancy grids

325 Two-dimensional occupancy grids are a fine representation for toy problems
 326 but they run into some obvious problems: since the occupancy grid is a “top view” of the world, we cannot represent non-trivial objects in it correctly
 327 (a large tree with a thin trunk eats up all the free space). We often desire a
 328 fundamentally three-dimensional representation of the physical world.
 329

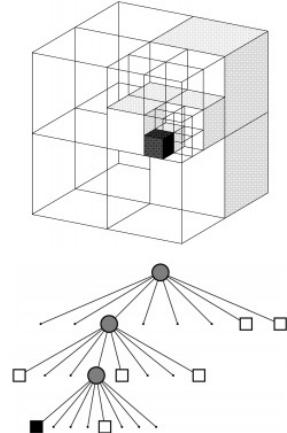


330

331 We could simply create cells in 3D space and our method for occupancy
 332 grid would work but this is no longer computationally cheap. For instance,
 333 if we want to build a map of Levine Hall (say 50 m × 50 m area and height
 334 of 25 m), a 3D grid map with a resolution of 5 cm × 5 cm × cm would have
 335 about 500 million cells (if we store a float in each cell this map will require
 336 about 2 GB memory). It would be cumbersome to carry around so many cells
 337 and update their probabilities after each sensor reading (a Velodyne gives data
 338 at about 30 Hz). More importantly, observe that most of the volume inside
 339 Levine is free space (inside of offices, inner courtyard etc.) so we do not really
 340 need fine resolution in those regions.

341 **Octrees** We would ideally have an occupancy grid whose resolution adapts
 342 with the kind of objects that are detected by the sensors. If nearby cells are
 343 empty we want to collapse them together to save on memory and computation,

344 on the other hand, if nearby cells are all occupied, we want to *refine* the
 345 resolution in that area so has to more accurately discern the shape of the
 346 underlying objects. Octrees are an efficient representation for 3D volumes.



347

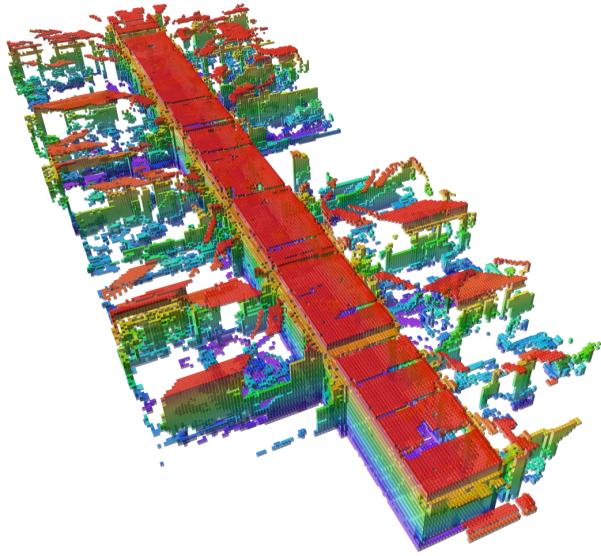
348 An octree is a hierarchical data structure that recursively sub-divides the 3D
 349 space into octants and allocates volumes as needed for a particular data point
 350 observed by a range sensor. It is analogous to a kd-tree. Imagine if the entire
 351 space in the above picture were empty (the tree only has a root node), and we
 352 receive a reading corresponding to the dark shaded region. An octree would
 353 sub-divide the space starting from the root (each node in the tree populates
 354 is the parent of its eight child octants) recursively until some pre-determined
 355 minimum resolution is reached. This leaf node is grid cell; notice how different
 356 cells in the octree have different resolutions. Occupancy probabilities of each
 357 leaf node are updated using the same formula as that of (4.23). A key point
 358 here is that octrees are designed for accurate sensors (LiDARs) where there is
 359 not much noise in the observations returned by the sensor (and thereby we do
 360 not refine unnecessary parts of the space)

361 Octrees are very efficient at storing large map, I expect you can store the
 362 entire campus of Penn in about a gigabyte. Ray tracing (following all the cells
 363 m_i in tree along the axis of the sensor in Figure 4.4) is harder in this case but
 364 there are efficient algorithms devised for this purpose. An example OctoMap
 365 (an occupancy map created using an Octree) of a building on the campus of
 366 the University of Freiburg is shown below.

❶ You can find LiDAR maps of the entire United States (taken from a plane) at <https://www.usgs.gov/core-science-systems/ngp/3dep>

367 4.5 Local Map

368 In this chapter, we primarily discussed occupancy grids of static environments
 369 as the robot moves around in the environment. The purpose of doing so is lo-
 370 calization, namely, finding the pose of the robot by comparing the observa-
 371 tions of the sensors with the map (think of the particle filter localization example
 372 in Chapter 3). In typical problems, we often maintain two kinds of maps, (i)
 373 a large occupancy grid for localization (say as big as city), and (ii) another
 374 smaller map, called the local map, that is used to maintain the locations of



375 objects (typically objects that can move) in the vicinity of the robot, say a 100
 376 m × 100 m area.

377 The local map is used for planning and control purposes, e.g., to check if
 378 the planned trajectory of the robot does not collide with any known obstacles.
 379 See an example of the local map at the 1:42 min mark at
 380 <https://www.youtube.com/watch?v=2va15BE-7lQ>. Some people also call the
 381 local map a “cost map” because occupied cells in the local map indicate a
 382 high collision cost of moving through that cell. The local map is typically
 383 constructed in the body frame and evolves as the robot moves around (objects
 384 appear in the front of the robot and are spawned in the local map and disappear
 385 from the map at the back as the robot moves forward).

You should think of the map (and especially the local map) as the filtering estimate of the locations of various objects in the vicinity of the robot computed on the basis of multiple observations received from the robot’s sensors.

386 4.6 Discussion

387 Occupancy grids are a very popular approach to represent the environment
 388 given the poses of the robot as it travels in this environment. We can also
 389 use occupancy grids to localize the robot in a future run (which is usually
 390 the purpose of creating them). Each cell in an occupancy grid stores
 391 the posterior probability of the cell being occupied on the basis of mul-
 392 tiple observations $\{y_1, \dots, y_k\}$ from respective poses $\{x_1, \dots, x_k\}$. This
 393 is a very efficient representation of the 3D world around us with the one
 394 caveat that each cell is updated independently of the others. But since one



Figure 4.6: The output of perception modules for a typical autonomous vehicle (taken from <https://www.youtube.com/watch?v=tiwVMrTLUWg>. The global occupancy grid is shown in gray (see the sides of the road). The local map is not shown in this picture but you can imagine that it has occupied voxels at all places where there are vehicles (purple boxes) and other stationary objects such as traffic light, nearby buildings etc. Typically, if we know that so and so voxel corresponds to a vehicle, we run an Extended Kalman Filter for that particular vehicle to estimate the voxels in the local map that it is likely to be in, in the next time-instant. The local map is a highly dynamic data structure that is rich in information necessary for planning trajectories of the robot.

395 gets a large amount of data from typical range sensors (a 64 beam Velodyne (<https://velodynelidar.com/products/hdl-64e>) returns about a 2 million
 396 points/sec and cheaper versions of this sensor will cost about \$100), this caveat
 397 does not hurt us much in practice. You can watch this talk
 398 (https://www.youtube.com/watch?v=V8JMwE_L5s0) by the head of Uber's
 399 autonomous driving group to get more perspective about localization and
 400 mapping.